



CLOUD COMPUTING APPLICATIONS

HBase Internals - Part 2

Roy Campbell & Reza Farivar

HRegion Assignment

- HBase Master keeps track of:
 - Set of live HRegion servers
 - Assignment of HRegions to HRegion servers
 - Unassigned HRegions
- Each HRegion is assigned to one HRegion server at a time
 - HRegion server maintains an exclusive lock on a file in ZooKeeper
 - HBase Master monitors HRegion servers and handles assignment
- Changes to HRegion structure
 - Table creation/deletion (HBase Master initiated)
 - HRegion merging (HBase Master initiated)
 - HRegion splitting (HRegion server initiated)

HRegion in Memory Representation

