



# **CLOUD COMPUTING APPLICATIONS**

Pregel - Part 1

Roy Campbell & Reza Farivar

# Introduction

- Infrastructure for graph processing – expensive to design
- Can we use MapReduce?
  - Inefficient because the graph state must be stored at each stage of the graph algorithm, and each computational stage will produce much communication between stages
- Single computer and library approach – not scalable
- Use existing shared memory parallel graph algorithms – no fault-tolerance

# Vertex-Oriented

- Based on BSP model
- Provides directed graph to Pregel
- Runs your computation at each vertex (processor)
- Repeats until every computation at each vertex votes to halt
- Pregel returns directed graph as a result

# Primitives

- Vertices – first class
- Edges – not first class
- Both vertices can be created and destroyed

# Pregel Organized via C++ API

- Supersteps  $S$
- Application code subclasses `Vertex`, writes a `Compute` method
- Can get/set `Vertex` value
- Can get/set outgoing edges values
- Can send/receive messages
- Reads messages sent to  $V$  in superstep  $S-1$ . Sends messages to other vertices that will be received at superstep  $S+1$ ; modifies state of  $V$  and its outgoing edges

# C++ API

- Message passing
- No guaranteed message delivery order
- Messages delivered exactly once
- Can send a message to any node
- If destination doesn't exist, user's function is called