



CLOUD COMPUTING APPLICATIONS

Trident: Exactly Once Processing

Roy Campbell & Reza Farivar

But What About State

- For most of storm, state storage is left up to you
- If your bolt goes down with 3 weeks of aggregated data that you have not stored anywhere, well too bad!

Enter Trident

- Provides exactly once semantics
- In trident, state is a first-class citizen, but the exact implementation of state is up to you
 - There are many prebuilt connectors to various NoSQL stores like HBase
- Provides a high level API (similar to cascading for Hadoop)

Trident Example

```
public class Split extends BaseFunction {  
  
    public void execute(TridentTuple tuple, TridentCollector collector) {  
        String sentence = tuple.getString(0);  
        for(String word: sentence.split(" ")) {  
            collector.emit(new Values(word));  
        }  
    }  
}
```

No Acking Required

Trident Example

```
TridentTopology topology =  
    new TridentTopology();  
TridentState wordCounts =  
    topology.newStream("spout1", spout)  
        .each(new Fields("sentence"), new Split(),  
            new Fields("word"))  
        .groupBy(new Fields("word"))  
        .persistentAggregate(new  
            MemoryMapState.Factory(), new Count(), new  
            Fields("count"))  
        .parallelismHint(6);
```

Aggregates values and stores them.

