

# Streaming Ecosystem

Reza Farivar

Capital One

*Reza.farivar@capitalone.com*

# Components of a streaming ecosystem

- Gather the data
  - Funnel
- Distributed Queue
- Real-Time Processing
- Semi-Real-Time Processing
- Real-time OLAP

# Step 1: Gather the Data

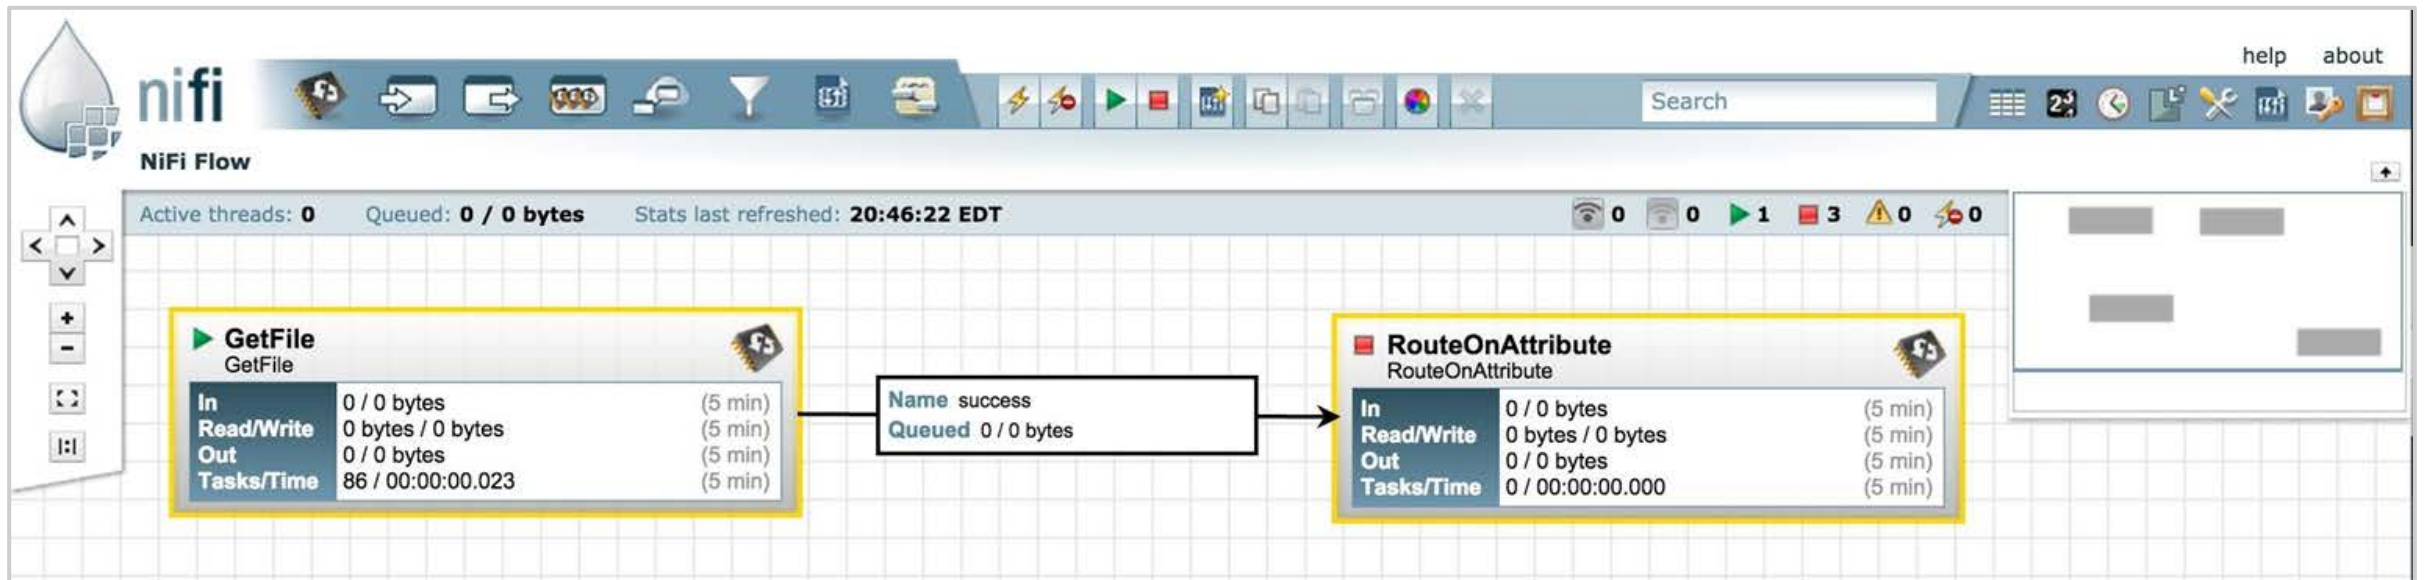
- Apache NiFi is a good distributed funnel
- Was made in NSA
  - Over 8 years of development
- Open sourced in 2014 and picked up by HortonWorks
- Great visual UI to design a data flow
- Has many many processor types in the box
- But not very good for heavy weight distributed processing
  - Same graph is executed on all the nodes

# NiFi Components

- FlowFile
  - Unit of data moving through the system
  - Content + Attributes (key/value pairs)
- Processor
  - Performs the work, can access FlowFiles
- Connection
  - Links between processors
  - Queues that can be dynamically prioritized
- Process Group
  - Set of processors and their connections
  - Receive data via input ports, send data via output ports

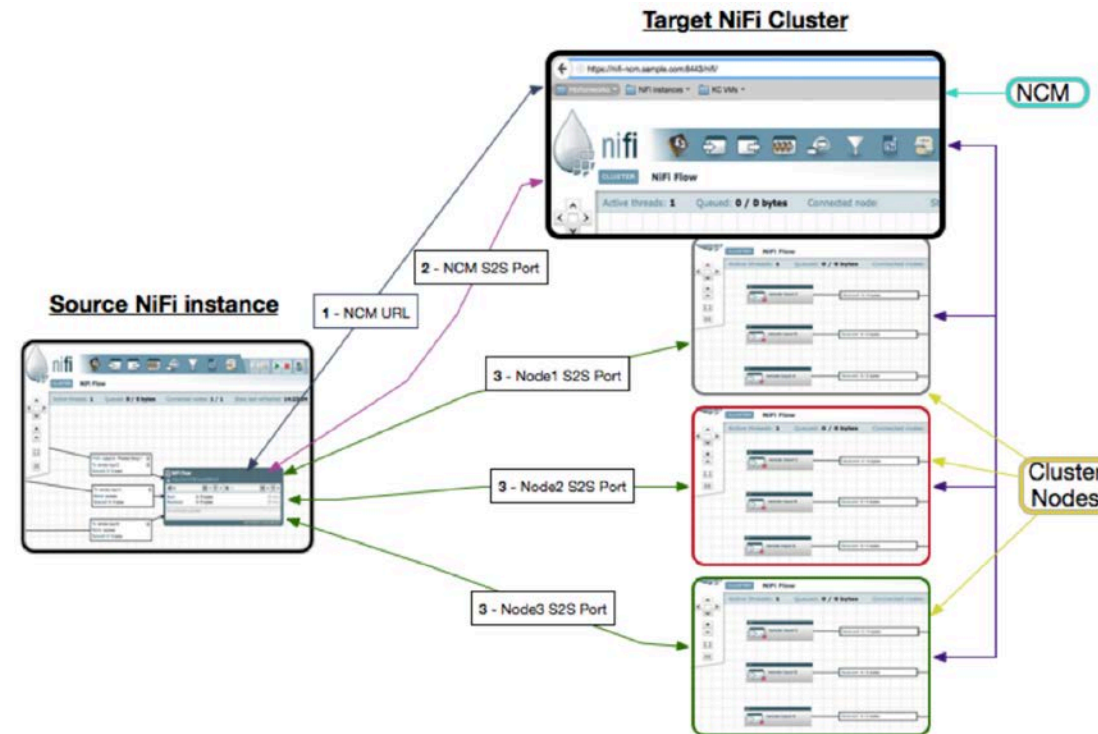
# NiFi GUI

- Drag and drop processors to build a flow
- Start, stop, and configure components in real time
- View errors and corresponding error messages
- View statistics and health of data flow
- Create templates of common processor & connections



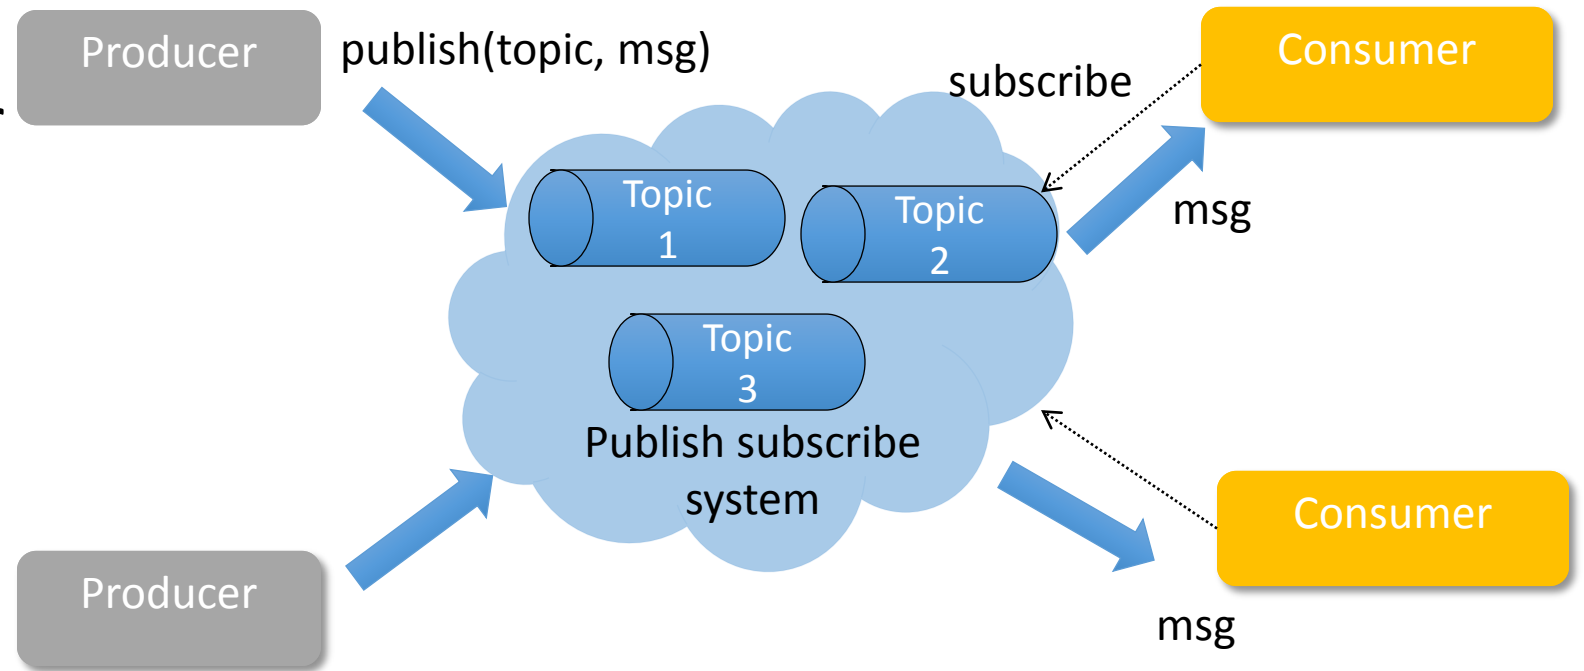
# NiFi Site-to-Site

- Site-to-site allows very easy pushing of data from one data center to another
- Makes it a great choice for distributed funnel




# Step 2: Distributed Queue

- Pub-sub model
- Kafka a very popular example



# Kafka Architecture

- Distributed, high-throughput, pub-sub messaging system
  - Fast, Scalable, Durable
- Main use cases:
  - log aggregation, real-time processing, monitoring, queueing
- Originally developed by LinkedIn
- Implemented in Scala/Java





# Kafka Manager

- There are some CLI tools

- `kafka-console-producer`

- `kafka-console-consumer`

- `Kafka-topics`

- `kafka-consumer-offset-checker`

- Some very new open-source projects for monitoring Kafka

- Kafka-manager by yahoo

- <https://github.com/yahoo/kafka-manager>

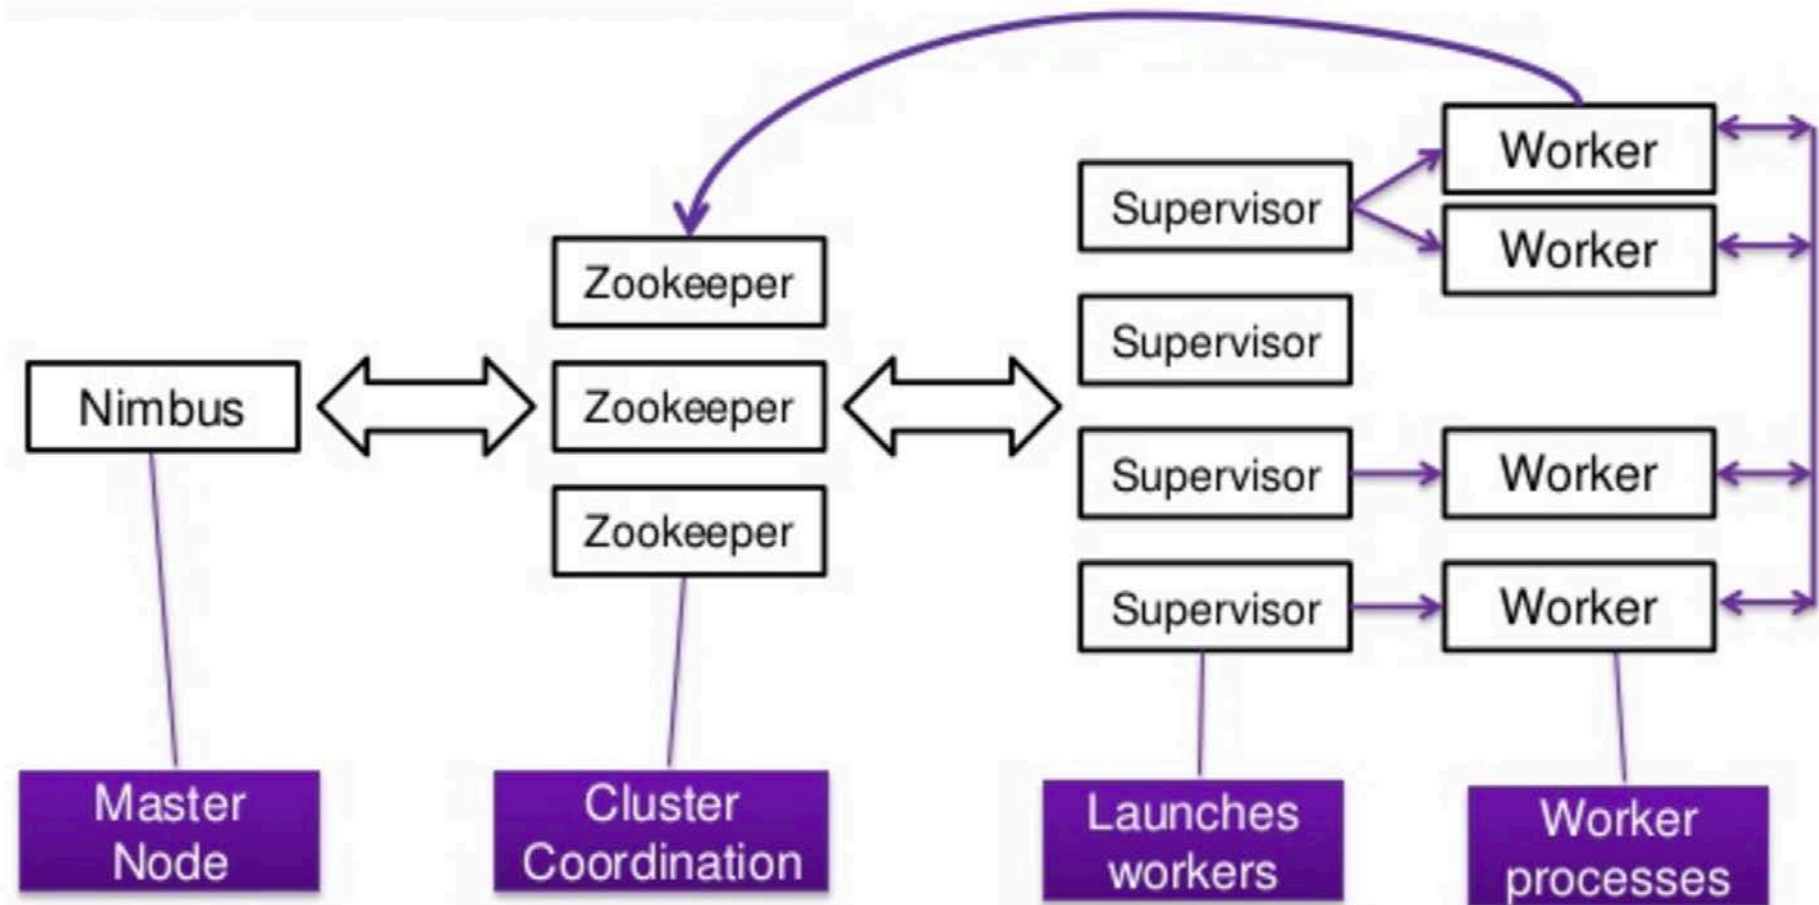
# Step 3: Distributed Processing

- Once data is in the Kafka message broker, we need to process it
- Filter
- Join
- Windowing
- Business logic
- Real-time requirements
  - Sub ms to 10 ms

# Storm

- Apache Storm
- Built in backtype, sold to Twitter
- Written in Clojure


# Storm Architecture



# Storm programming

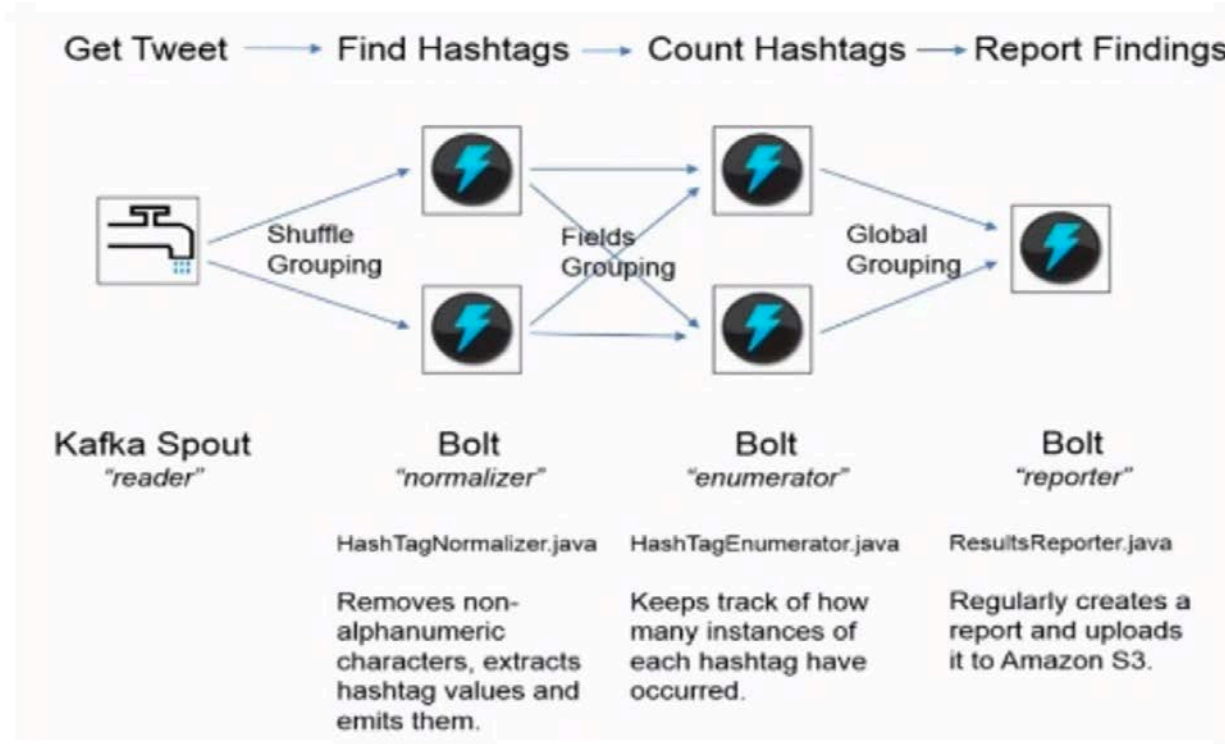
- Topology
  - Spouts
  - Bolts
  - Tuples
  - Streams
- topologyBuilder API

```
TopologyBuilder builder = new TopologyBuilder();  
builder.setSpout("words", new TestWordSpout(), 10);  
builder.setBolt("exclaim1", new ExclamationBolt(), 3)  
    .shuffleGrouping("words");  
builder.setBolt("exclaim2", new ExclamationBolt(), 2)  
    .shuffleGrouping("exclaim1");
```



# Example topology

- Storm is great for non-trivial large scale processing
- Mature enterprise level features, including multitenancy and security
- Work on resource aware scheduling



## Step 5: Micro batch processing / SQL / ML

- Instead of real-time event-by event processing, we can do micro batch
- Reduce overheads
- Fault tolerance → Kappa architecture
- High latency


# Spark

- Spark was a project out of Berkeley from 2010
- Has become very popular
- Most contributed open source project in big-data domain
- RDD: Resilient Distributed Data Set



# Spark Streaming

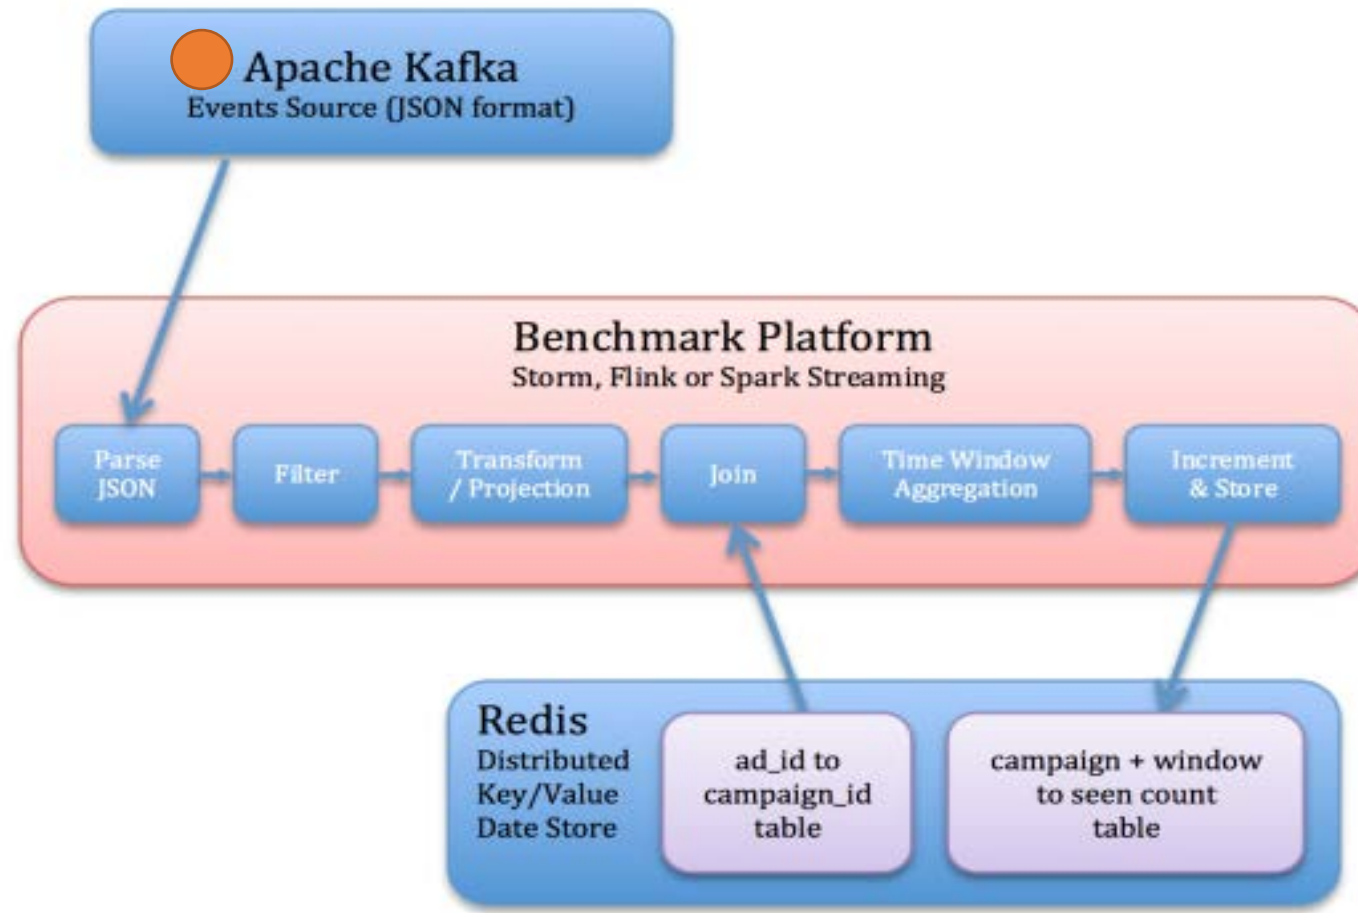
- Window a bit of data
- Run a batch
- Repeat



# Spark ML, Graph, etc.


- Advantage of Spark Streaming:
  - Rich ecosystem of big data tools
  - Spark SQL
  - Spark ML
  - Spark GraphX
  - SparkR
- Disadvantage:
  - Not really streaming

# Benchmark: ETL pipeline



# Three-way Comparison

- Flink and Storm have similar linear performance profiles
  - These two systems process an incoming event as it becomes available
- Spark Streaming has much higher latency, but is expected to handle higher throughputs
  - System behaves in a stepwise function, a direct result from its micro-batching nature



# Side note: in-memory key-value store

- Redis
- Cassandra

# Step 6: OLAP (Online Analytical Processing)


- Business Intelligence
- Multidimensional data analytics
- Analyze multidimensional data interactively
- Basic Operations
  - Consolidation (roll-up, aggregation in dimensions)
  - Drill-down (filter)
  - Slicing and dicing (Look at the data from different viewpoints)

# Druid

- Developed in Metamarkets in 2011
  - RDBMs: Too slow
  - NoSQL key value store: fast, but exponential memory space, precompute very slow
- Gaining in popularity
- Open Source (Apache license) in late 2012
- OLAP queries
- Column oriented
- Sub second query time (Avg query time 0.5 seconds)
- Real-time streaming ingestion
- Scalable

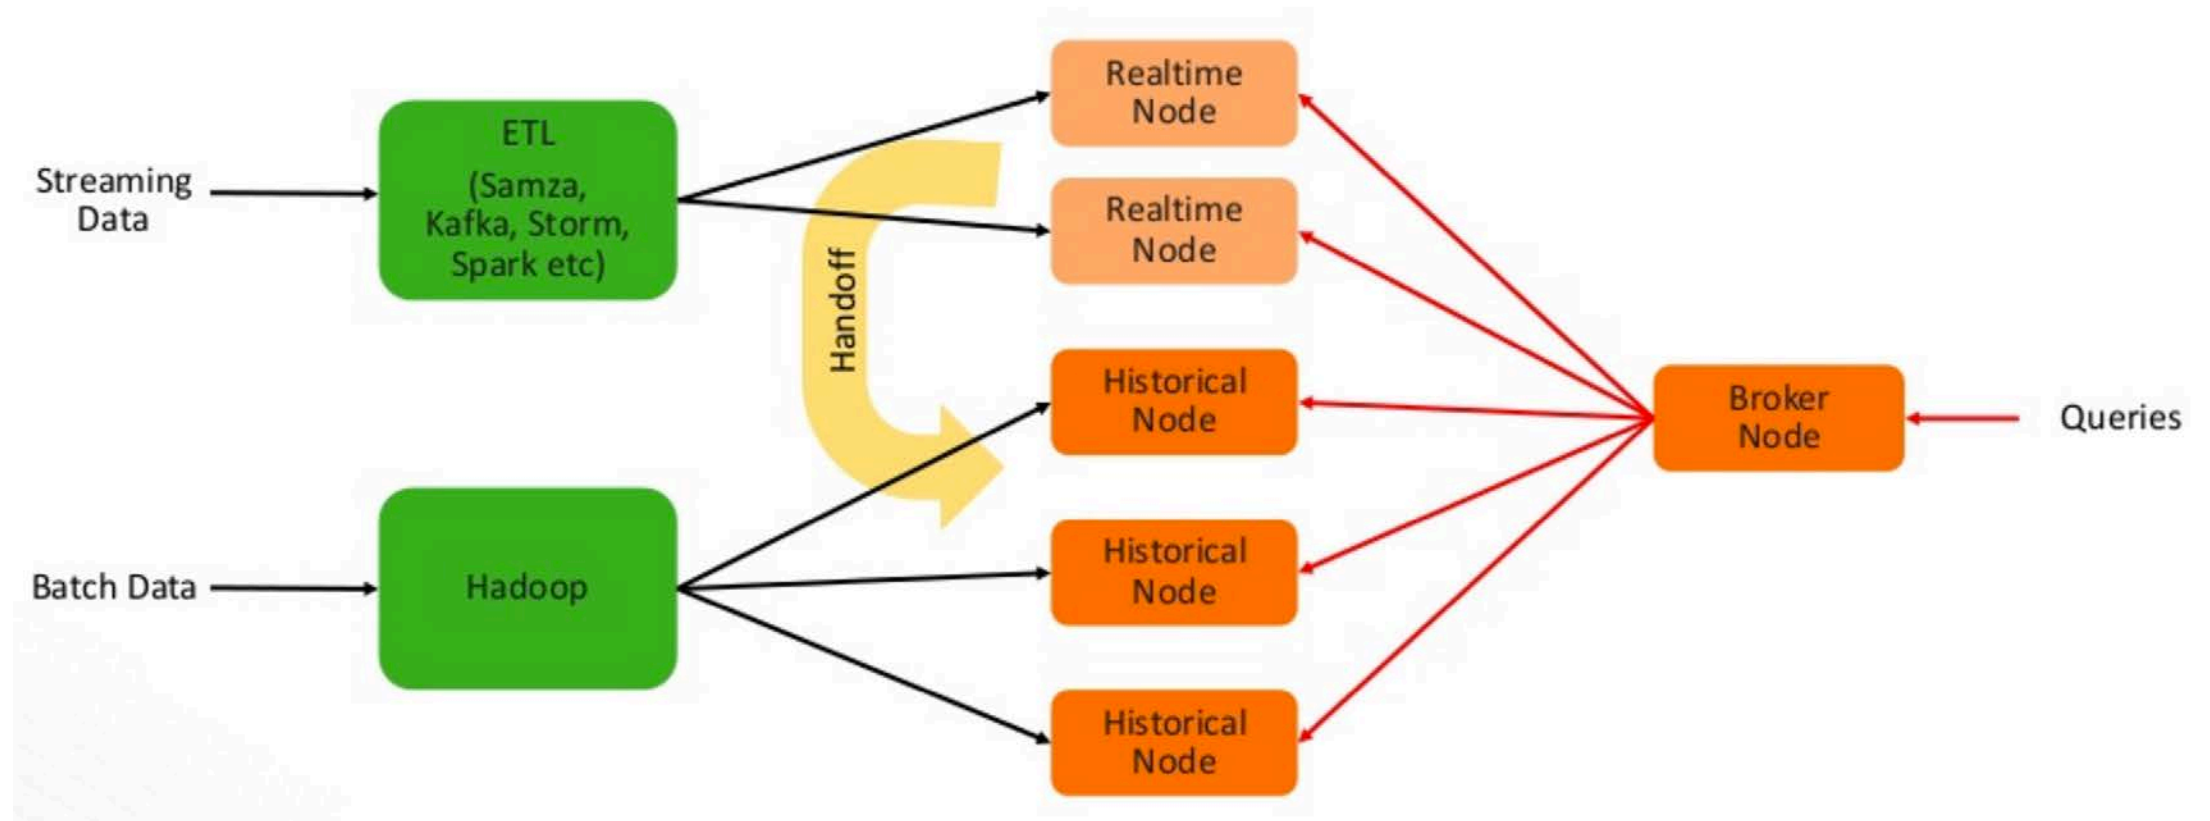
# Druid

- Arbitrary slice and dive of data





# Druid Architecture



# Druid Bitmap Index

- This is one of the reasons Druid is so fast
- Dictionary encoding
- Bitmap Index
- Compression ratio: 1 bit per record
- Logical AND/OR of a few thousand numbers for a query → lightning fast queries


timestamp	page	language	city	country	...	added	deleted
2011-01-01T00:01:35Z	Justin Bieber	en	SF	USA		10	65
2011-01-01T00:03:63Z	Justin Bieber	en	SF	USA		15	62
2011-01-01T00:04:51Z	Justin Bieber	en	SF	USA		32	45
2011-01-01T01:00:00Z	Ke\$ha	en	Calgary	CA		17	87
2011-01-01T02:00:00Z	Ke\$ha	en	Calgary	CA		43	99
2011-01-01T02:00:00Z	Ke\$ha	en	Calgary	CA		12	53
...							

▶ Justin Bieber -> [0, 1, 2] -> [111000]  
▶ Ke\$ha -> [3, 4, 5] -> [000111]

# Step 7: BI

- Pivot
  - web-based exploratory visualization UI for Druid
  - Easily filter, split, visualize, etc.
- Tableau and SQL not natively supported 😞
  - But wait!

# Pivot




# Druid and Spark

- Druid's native API is JSON
- No Tableau, SQL support
- But there is hope!

<https://github.com/SparklineData/spark-druid-olap>


- Connect Druid to Tableau through Spark

```
CREATE TEMPORARY TABLE orderLineItemPartSupplier  
USING org.sparklinedata.druid  
OPTIONS (sourceDataframe "orderLineItemPartSupplierBase",  
timeDimensionColumn "l_shipdate",  
druidDatasource "tpch",  
druidHost "localhost",  
druidPort "8082",  
columnMapping '{ "l_quantity" : "sum l_quantity",  
                "ps_availqty" : "sum ps_availqty"  
            })
```



# Why Druid and Spark together?

- Spark is great as a general engine
- Everything and the kitchen sink
- Queries can take a long time
  - Still much faster than Hive on Yarn
- Druid is optimized for Column based time-series queries



# Questions?

Email: [reza.farivar@capitalone.com](mailto:reza.farivar@capitalone.com)